

# Linux Distro Tools for Building Container Images

Nisha Kumar (@nishakmr)

Joshua Lock (@hi\_joshuagl)

Open Source Software Engineers / OSTC

@vmwopensource

November 2019

# Agenda

## Motivations and context

Motivations for exploring other container image build tools  
Context and level setting

## Investigation

Linux distros and tools  
Experiments and analysis

## Summary

Takeaways  
Actionable learning

# Motivations and context

Why are we looking at distro tools?

# Motivations



Compliance



Security



Discipline

# Level setting

Containers ~= packaging format



# Desirable properties of a container image

...or package format

Repeatable

At any given time, we can reliably create an equivalent image from source

Identifiable

We can reason about their contents

Recent

Contents are not old and vulnerable

# Common practices which violate the desirable properties

Performing non-deterministic operations (RUN)

Encoding build-time dependent state in the image (apt-get update)

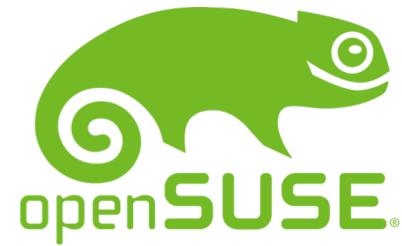
Inserting untraceable files into the filesystem (ADD/COPY)

Using an old base OS (FROM)



# These problems have been solved before

Why reinvent the wheel?



# Tool usability continuum

...or DevOps



# Investigation

Using distribution tools to build container images

# Linux distributions and tools

## Choosing is hard

### Debian and DebOS

- Debian – the universal OS
- DebOS – Debian image generator created for Embedded Linux use-cases

### Buildroot

- Embedded Linux image creator

### Yocto Project

- Embedded Linux distribution builder

### Guix SD

- Functional distro built on Guix package manager, inspired by Nix
- Scheme extensions provide DSLs for packaging and configuration

# Distribution tool showdown

Trying to use the tools and compare the outputs

## Compare

- Ease of use
- Output image size
- Engineering effort
- Quality/presence of SBoM



[This Photo](#) by Unknown Author is licensed under [CC BY-SA-NC](#)

# Experiment 1: base OS for Go applications

## Base OS with Go toolchain

Goal: create a container image with Go toolchain

DebOS:

- Ease of use
- Image size
- Engineering effort
- SBoM



# Experiment 1: base OS for Go applications

## Base OS with Go toolchain

Goal: create a container image with Go toolchain

### Buildroot:

- Ease of use
- Image size
- Engineering effort
- SBoM



??



# Experiment 1: base OS for Go applications

## Base OS with Go toolchain

Goal: create a container image with Go toolchain

### Yocto Project:

- Ease of use
- Image size
- Engineering effort
- SBoM



# Experiment 1: base OS for Go applications

## Base OS with Go toolchain

Goal: create a container image with Go toolchain

### Guix:

- Ease of use
- Image size
- Engineering effort
- SBoM



# Status Quo -> Just Enough

Current

Base OS

+

App SDK

+

App

Ideal

App

+

Dependencies

# Packaging



---

*Package:* describe a piece of software, its dependencies and how to build/install it

*Package manager:* handles installation and removal of packages (and dependencies)

Can we leverage the  
package manager to install  
*only* what's *required* into an  
image?

# Experiment 2: Just enough OS (JeOS)

Most container workloads don't need a full OS root filesystem

Goal: create  
minimal  
container image  
with only the  
containerised  
app and its  
required  
dependencies.



# Packaging for our target distro tools

My application isn't packaged

Each distro ecosystem has different packaging and metadata formats and tools to process them.

Tools exist to automate packaging for many Linux distributions and tools

- Their support for cloud native languages and frameworks (Go, Bazel, Java, etc) is patchy to non-existent

Automatic packaging tools:

Debian — dh-make / dh-make-golang / gem2deb / pypi2deb

Buildroot — docs for packaging for each ecosystem

Yocto Project — recipetool

Guix — guix import



# What if we *could* use the native package?

Eg: Postgresql

Current



Base OS

+



Dependencies

+



App

Ideal



App

+

Dependencies

# Summary

Closing remarks

# Takeaways

This is only a partially solved problem

Tools build for different hardware now need to be adapted for different apps

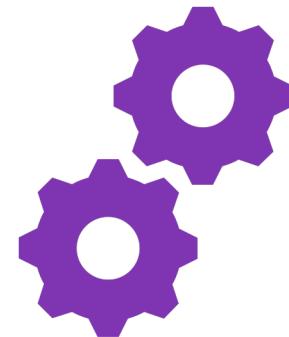
Projects are working on making the jump to cloud

Compelling reasons to help them (Open Source!)

- Documentation

- Tests

- Use cases



# Actionable learning

What can we do today for my Dockerfile built containers?

Introspect your containers: Dive, Container-diff, Tern

Get to halfway

- Fairly straightforward to use debos

- Make your Base OS + SDK with debos

- Use this Base OS for containerization

- Glue code needed for Guix

- Make the appropriate rootfs from gnu/store

- Upstream above improvements to make it seamless

# Resources

DebOS: <https://github.com/go-debos/debos>

Debian packaging:

<https://wiki.debian.org/Packaging/Intro?action=show&redirect=IntroDebianPackaging>

Buildroot user manual: <https://buildroot.org/downloads/manual/manual.html>

Guix user manual: [https://guix.gnu.org/manual/en/html\\_node/](https://guix.gnu.org/manual/en/html_node/)

Yocto Project Mega Manual:

<https://www.yoctoproject.org/docs/3.0/mega-manual/mega-manual.html>

Our Lab Notebook: <https://github.com/nishakm/distro-tools-for-containers>



Joshua: @hi\_joshuagl

Nisha: @nishakmr

**Thank You**

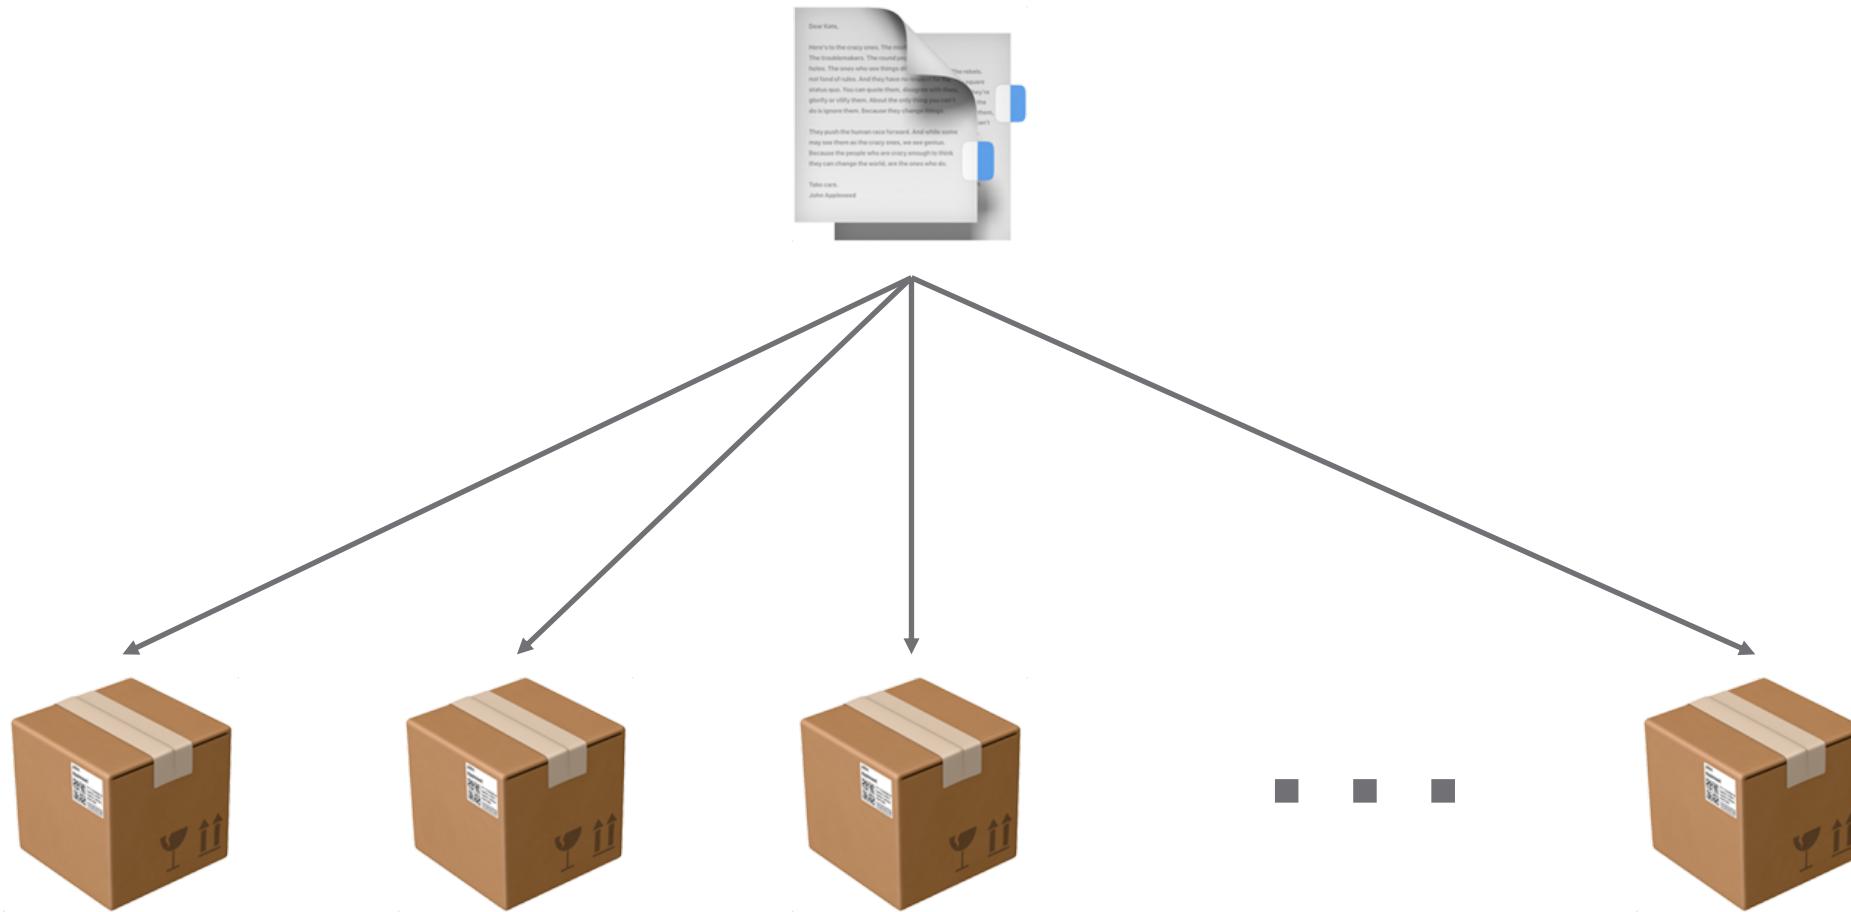
# Sidebar: Image layers

Cache invalidation, image size, etc.

IMAGE	CREATED	CREATED BY	SIZE	COMMENT
a8c1ae078d5b	10 seconds ago	/bin/sh -c #(nop) CMD ["/bin/sh" "-c" "yarn...	0B	
9d2d1b910ebc	10 seconds ago	/bin/sh -c apt-get remove -y git python3-pip...	1.91MB	
085a5d2da5f4	13 seconds ago	/bin/sh -c tar -xvzf gosec_2.0.0_linux_amd64...	8.85MB	
f8290ae97e6a	5 days ago	/bin/sh -c curl -OL https://github.com/secur...	4.72MB	
d46051df15a7	5 days ago	/bin/sh -c yarnpkg install	272MB	
bddd59b50ced	5 days ago	/bin/sh -c #(nop) WORKDIR /precaution	0B	
396a54161949	5 days ago	/bin/sh -c git clone https://github.com/vmwa...	2.43MB	
8ab12298d2da	5 days ago	/bin/sh -c apt-get update && apt-get upgrade...	583MB	
85c4fd36a543	6 weeks ago	/bin/sh -c #(nop) CMD ["bash"]	0B	
<missing>	6 weeks ago	/bin/sh -c #(nop) ADD file:99bf629976cd3d79c...	114MB	

# Sidebar: OCIv1 Image Format

Tarballs all the way down



# Sidebar: OCIv1 Image Format

Tarballs all the way down



=



=



=



# Experiment 2a: packaging Fluentd

## Unified logging layer

Why fluentd? ([fluent.org](https://fluent.org))

- Cloud native application (we all ❤️ logs)
- Ruby - dynamic languages require more diligence

Debian: 😞

- gem2deb requires dependencies be satisfiable by apt
  - Ergo; need to package dependencies first

Buildroot: 😱

- No auto-packaging tools exist for buildroot

Yocto Project: 😟

- recipetool doesn't support Ruby projects

Guix: 😢

- + guix import has a --recursive option, to also handle dependencies
- Still couldn't build a working package